

AgentOps: Monitor and govern AI agents with IBM Telemetry by using watsonx Orchestrate

Introduction

As AI agents become more sophisticated and autonomous, understanding their behavior, performance and decision-making processes is critical for ensuring reliability and governance. [AgentOps](#), the practice of monitoring, observing and managing AI agents in production, provides the visibility needed to build trustworthy agentic AI systems.

This tutorial provides a step-by-step guide to setting up and using IBM Telemetry with watsonx Orchestrate® Developer Edition to monitor and govern AI agents. You'll learn how to enable observability for AI agents and analyze their behavior in depth, from individual LLM calls to complete multistep workflows.

By the end of this tutorial, you'll be able to:

- Install and configure watsonx Developer Edition locally
- Enable IBM Telemetry for comprehensive agent observability
- Import and test a preconfigured AI agent with external tool integration
- Analyze agent behavior through detailed traces, tasks, spans and workflows
- Debug issues and optimize agent performance by using advanced analytics

What is IBM Telemetry?

IBM Telemetry is the native observability framework of watsonx Orchestrate that captures detailed information about how your AI agents execute requests. It records every step of the agent lifecycle, from routing decisions and prompt construction to LLM invocations and tools calls, providing complete visibility to agent behavior.

With IBM Telemetry, you can track performance metrics, monitor LLM cost, identify errors and ensure that your agents are operating as intended. IBM Telemetry provides enterprise-grade observability designed for production environments and AI systems at scale.

Prerequisites

System requirements

Before you begin, ensure that the following prerequisites are installed and configured on your system:

- Python 3.8+ (Check with `python --version`)
- 16 GB RAM minimum
- The watsonx Orchestrate Developer Edition through watsonx Orchestrate ADK

This guide includes installation steps for the ADK.

Authorization requirements

- **watsonx Orchestrate Account** [Free 30-day trial available](#) if you don't have an account.
- **watsonx Orchestrate API Key**, generated from your IBM Cloud® dashboard

Authorization steps are provided later in this guide.

Steps

Step 1. Clone the GitHub Repository

To get started, clone the GitHub repository by using <https://github.com/IBM/ibmdotcom-tutorials.git> as the HTTPS URL. For detailed steps on how to clone a repository, refer to the [GitHub documentation](#).

Open the repository in your preferred integrated development environment (IDE) (for example, [Visual Studio Code](#)) and locate this tutorial's project folder: `wxo-agentops`. This directory is where you'll be working in as you follow along.

Step 2. Install the watsonx Orchestrate ADK

The IBM watsonx orchestrate Agent Development Kit (ADK) is a CLI tool that simplifies the installation, configuration and management of the watsonx Orchestrate Developer Edition.

To use the ADK, you must connect it to an existing watsonx Orchestrate environment. If you don't have a watsonx Orchestrate account yet, you can [sign up for a free 30-day trial](#). If you already have an account, you can use it to provide the environment credentials needed by the ADK.

These steps will guide you through installation by using a Python virtual environment, which is the recommended approach for keeping dependencies isolated. For alternative installation methods and detailed instructions, see the Getting started with the ADK documentation.

2a. Create your virtual environment

Create a new Python virtual environment in your project directory:

`python -m venv .venv`

Powered by  Granite

This step creates a `.venv` folder containing an isolated Python environment.

2b. Activate your virtual environment

The activation command differs depending on your operating system.

macOS and Linux

`source ./venv/bin/activate`

Powered by  Granite

Windows

.\venv\Scripts\activate

Powered by  Granite

Once activated, your terminal prompt should change to indicate you're working inside the virtual environment (typically showing (.venv) at the beginning of the prompt).

2c. Install the watsonx Orchestrate ADK

With your virtual environment activated, install the ADK by using pip:

pip install ibm-watsonx-orchestrate

Powered by  Granite

This command downloads and installs the ADK along with all its dependencies. The installation can take a few minutes to complete.

Note: If you have an earlier version of the ADK installed (>2.0), run pip install --upgrade ibm-watsonx-orchestrate . You might also have to run the troubleshooting steps in step 4b.

Step 3. Configure your environment

The ADK uses a .env file to authenticate your user credentials and configure the watsonx Orchestrate Developer Edition. The environment variables that you need depend on your chosen authentication method. This tutorial uses the watsonx Orchestrate account method, which is the most straightforward approach for getting started.

For alternative authentication methods and detailed configuration instructions, see the [configuring your environment file documentation](#).

Step 3a. Create your .env file

Step 3b. Configure required fields

Open the .env file in your text editor and configure the following two essential fields:

- **WO_INSTANCE :** This URL is your watsonx Orchestrate instance. You can find this information by logging in to your [watsonx Orchestrate account](#) and navigating to your instance details. Click your profile icon > Settings, then select the **API details** tab. For detailed instructions on getting started with the API, check out the [watsonx Orchestrate documentation](#).

The URL follows this format:

WO_INSTANCE=https://api.us-south.watson-orchestrate.cloud.ibm.com/instances/<your-instance-id>

Powered by  Granite

Copy and paste your service instance URL to replace the template value in your .env file. The region (for example, us-south depends on your geographical location).

- **WO_API_KEY** : This key is your watsonx Orchestrate application programming interface (API) key, which authenticates your connection to IBM Cloud services. You can generate or retrieve this key from your IBM Cloud account dashboard. Replace <your-api-key> with your actual API key. For step-by-step instructions on generating an API key, see the [getting started documentation](#).

WO_API_KEY=<your-api-key>

Powered by  Granite

Keep your API key secure and never commit it to version control. The .env file should already be included in your .gitignore to prevent accidental exposure.

Step 4. Install the watsonx Orchestrate server and enable IBM Telemetry

Now, you're ready to install the [watsonx Orchestrate Developer Edition](#), which will run a local instance of the watsonx Orchestrate server on your machine. This step also enables IBM Telemetry, giving you immediate access to observability features.

Understanding the installation command

The ADK provides a single command that handles the entire installation process:

orchestrate server start -e <path-.env-file> --with-ibm-telemetry

Powered by  Granite

Let's break down what this command does:

- **orchestrate server start** : Initializes and starts the watsonx Orchestrate Developer Edition server
- **-e <path-.env-file>** : Points to your configuration file containing credentials
- **--with-ibm-telemetry** : Enables the native observability framework of IBM Telemetry

4a. Run the installation

Execute the command from your wxo-agentops directory:

The following command starts the watsonx Orchestrate Developer Edition server by initializing the server environment: `orchestrate server start -e <path-.env-file>`. Adding the `--with-ibm-telemetry` flag enables IBM Telemetry, its native observability framework.

Run this command to install watsonx Orchestrate server with IBM Telemetry:

`orchestrate server start -e .env --with-ibm-telemetry`

Powered by  Granite

This command creates internal containers managed by the ADK for:

- The watsonx Orchestrate server

- PostgreSQL and Redis databases
- IBM Telemetry services
- Supporting dependencies

The ADK automatically configures a virtual network that allows these containers to communicate with each other on <http://localhost:3000>.

Step 4b. Verify successful installation

The installation process can take several minutes, especially on the first run as the necessary images are downloaded. A successful installation produces output similar to this example:

[INFO] - Waiting for orchestrate server to be fully initialized and ready... [INFO] - Orchestrate services initialized successfully [INFO] - local tenant found [INFO] - You can run `orchestrate env activate local` to set your environment or `orchestrate chat start` to start the UI service and begin chatting.

Powered by  Granite

If you see this message, congratulations! Your local watsonx Orchestrate environment with IBM Telemetry is now running.

Troubleshooting installation

If the installation fails or hangs, try the following steps:

1. Reset the server:

orchestrate server reset

Powered by  Granite

This command stops and removes all containers created for watsonx Orchestrate, giving you a clean slate.

2. Restart the installation:

After resetting, run the start command again:

orchestrate server start -e .env --with-ibm-telemetry

Powered by  Granite

3. Check the server logs container status:

You can view service logs for the Orchestrate server to check for warnings or errors:

orchestrate server logs

Powered by  Granite

If the preceding steps do not work, reset the server and completely removing the server environment: orchestrate server purge and reinstall.

Step 5. Activate your local environment and start the service

With the watsonx Orchestrate server successfully installed, you now need to activate your local environment and launch the chat interface where you'll interact with your AI agents.

Activate the local watsonx Orchestrate environment

The watsonx Orchestrate ADK supports multiple environments (local, development, production, and so on.). You need to explicitly activate the local environment you created:

orchestrate env activate local

Powered by  Granite

You should receive confirmation that the environment is active:

[INFO] - local tenant found [INFO] - Environment 'local' is now active

Powered by  Granite

This sets the local environment as your default context for all subsequent ADK commands. Any agents, tools or configurations you work with will now target this local instance.

Launch the watsonx Orchestrate chat interface

Start the watsonx Orchestrate chat UI service with the following command:

orchestrate chat start

Powered by  Granite

This command initializes the web-based chat interface and automatically opens it in your default browser. You should see output similar to:

[INFO] - Chat UI Service started successfully. [INFO] - Waiting for UI component to be initialized...

[INFO] - Opening chat interface at http://localhost:3000/chat-lite

Powered by  Granite

The chat interface provides an easy way to interact with your AI agents. If the browser doesn't open automatically, you can manually navigate to http://localhost:3000/chat-lite.

Verify that the interface is running

Once the chat interface loads, you should see a clean chat window ready for interaction. At this stage, you haven't imported any agents yet, so the interface will be mostly empty. That result is expected, you'll add your first agent in the next step.

Step 6. Import a weather agent and tool to test IBM Telemetry

Now that your environment is set up, it's time to import a preconfigured AI agent that demonstrates the monitoring capabilities of IBM Telemetry. This weather agent uses an external API tool to fetch real-time weather data, giving you a practical example to observe and analyze.

Why start with a weather agent?

The weather agent is an ideal starting point because it:

- **Demonstrates tool usage:** Shows how agents call external APIs

- **Provides clear, observable behavior:** Each request follows a predictable pattern
- **Generates meaningful telemetry data:** Produces rich traces that you can analyze in IBM Telemetry
- **Includes error scenarios:** Helps you understand how telemetry handles failures
- **Illustrates automation:** Removes manual data lookup through agent actions

Step 6a. Navigate to the weather agent directory

From your project root (wxo-agentops), navigate to the Weather Agent folder:

```
cd weather_agent
```

Powered by  Granite

This directory contains two YAML configuration files:

- `get_weather.yaml` : Defines the weather API tool
- `weather_agent.yaml` : Defines the agent that uses this tool

Step 6b. Import the weather tool

Tools are reusable capabilities that agents can invoke to perform specific actions. Import the `get_weather` tools first:

```
orchestrate tools import -f get_weather.yaml --kind openapi
```

Powered by  Granite

The `--kind openapi` flag indicates that this tool uses an OpenAPI specification to define its interface. You should see confirmation that the tool was successfully imported.

Step 6c. Import the weather agent

Now import the agent that will use this tool:

```
orchestrate agents import -f weather_agent.yaml
```

Powered by  Granite

This command registers the Weather Agent in your local Watsonx Orchestrate environment. The agent is preconfigured with:

- Instructions on how to interpret weather data
- Permission to call the `get_weather` tool
- Fallback behavior for invalid locations

Step 6d. Activate the agent in the chat interface

Return to your browser where the chat interface is running. You might need to refresh the page to see the newly imported agent.

Click the agent dropdown menu (typically located at the top of the chat interface) and select **Weather_Agent** from the list

Test the agent

With the Weather Agent selected, try asking some questions to generate telemetry data:

Example queries:

- “What’s the weather like in New York City?”
- “Can you tell me the current temperature in London?”
- “What is the weather in Tokyo?”
- “Is it raining in Seattle right now?”

The agent will process each request by:

1. Understanding your query
2. Extracting the location
3. Calling the get_weather tool with appropriate coordinates
4. Interpreting the weather data
5. Responding in natural language

What's happening behind the scenes?

Every interaction you have with the Weather Agent is being captured by IBM Telemetry. The system is recording:

- The full conversation context
- Every LLM invocation and the tokens used
- Tool calls with their inputs and outputs
- Routing decisions and workflow steps
- Execution times and performance metrics
- Any errors or exceptions that occur
- Interactions with external providers and their response times

In the next step, you’ll explore this telemetry data in detail to understand exactly how your agent behaves.

Step 7. Analyze the agent's behavior in IBM Telemetry

Now comes the most powerful part of this tutorial: using IBM Telemetry to gain deep visibility into your agent’s behavior. IBM Telemetry provides multiple views and analytics tools that let you understand every aspect of how your agent processes requests.

Step 7a. Access the IBM Telemetry interface

Open your browser and navigate to <https://localhost:8765/?serviceName=wxo-server>. The interface provides session replays that let you revisit past agent interactions for analysis.

Note: The URL uses https but because this space is a local development environment, your browser might show a security warning about a self-signed certificate. This message is expected and it is safe to proceed past in your local environment.

Step 7b. Log in to IBM Telemetry

When the login screen appears, enter any name (to identify your local session) and click **Login**.

You'll be taken to the main IBM Telemetry dashboard.

Step 7c. Navigate to the Trace and Group Selection view

The dashboard shows a list of recent traces, each representing a single user interaction with an agent. Click the first trace in the **Trace and Group Selection** panel to view detailed analytics about your most recent chat with the Weather Agent.

This step takes you to the **Agent Analytics** screen, which serves as the central hub for understanding agent behavior.

Understanding the Agent Analytics screen

The Agent Analytics screen provides an overview of the selected trace, including:

- **Summary statistics:** Total execution time, token usage, cost estimates and benchmarking data
- **Agent information:** Which agent handled the request
- **User query:** The original question asked
- **Response preview:** The agent's final answer
- **Status indicators:** Success, warnings or errors

This high-level view gives you immediate insight into whether the agent performed as expected and how efficiently it operated.

In-depth review: Observe agent tasks

The **Tasks** section is where you'll spend most of your time analyzing agent behavior. It provides a visual, step-by-step timeline of everything the agent did during a request (every LLM call, tool invocation, routing decision and output generation).

Tasks are organized hierarchically to reflect how the agent actually executed the workflow, making it easy to understand the sequence of operations and their relationships.

Breaking down the Weather Agent task workflow

Let's examine the standard execution path for a watsonx Orchestrate agent request. Your Weather Agent trace should show a structure similar to this example:

```
0:_ROOT 0.0:agent_style_router # Routes the request
0.1:agent # Prepares prompt + logic
0.1.0:WatsonxChatModel.chat # LLM processes the request
0.2:answer # Sends final answer to user
Powered by  Granite
```

This workflow shows the entire lifecycle of a single user query. Here's what each task represents:

- **0:_ROOT :** The top-level span that contains all child tasks. Think of this folder as the one that holds the entire agent execution. It defines the start and end time of the complete trace, from the when the request enters the system until the final response is delivered.

This approach matters because the root task's duration tells you the total latency the user experienced. If the number is too high, you can analyze its child tasks to identify bottlenecks.

- 0.0:agent_style_router : The routing task determines which agent should handle the message and classifies the request into a handling style. The router analyzes the incoming request and decides whether it requires conversational handling, tool-driven execution, retrieval-augmented generation (RAG) or multi-agent orchestration.

The router ensures that the correct downstream logic is invoked. If requests are being misrouted, this is where you'd identify the problem.

- 0.1:agent : The main agent execution context that orchestrates the entire request. This task assembles the prompt from system interactions, conversation history and tool responses. It applies orchestration rules and policies and prepares inputs for the LLM. This task determines what type of LLM call to make.

This step is where the “intelligence” of orchestration happens. The agent task ensures that the LLM receives all the context it needs to make informed decisions.

- 0.1.0:WatsonxChatModel.chat : The actual call to the LLM where it receives the full prompt and decides whether to call a tool, ask for clarification or produce a direct answer. It generates the response, in either text or structured tool calls.

This step is the “thinking” step where the model processes information and makes decisions. Token usage, latency and quality issues all stem from this task. If your agent is slow or expensive, this step is usually the primary contributor.

- 0.2:answer : The final step in the chain takes the LLM’s output and formats it for delivery. This task converts the LLM’s raw output into the final answer format and applies any post-processing or formatting rules. Finally, it delivers the response back to the chat interface.

This task ensures that the user receives a properly formatted response. If answers are being truncated or improperly formatted, this step is where you'd investigate.

Task workflow summary

To summarize the complete workflow:

1. Router decides how to handle the request
2. Agent prepares the context and orchestration logic
3. LLM generates the response or tool calls
4. Answer formats and returns the final output

All of this workflow is wrapped under the ROOT request container, giving you a complete picture of the agent’s execution from start to finish. This level of observability is essential for MLOps and DevOps teams managing agent operations and complex pipelines at scale.

Understanding task attributes

Each task in the hierarchy contains three categories of attributes that provide detailed metadata about what the task consumed and produced:

1. **Input attributes:** Show everything the task received before execution: messages, tool responses, system instructions, the internal state.

Example: For the WatsonxChatModel.chat task, input attributes would include the fully assembled prompt with system instructions, conversation history and any tool results that need to be interpreted.

2. **Output attributes:** Show what the task produced, including: LLM completions, tool calls and decisions.

Example: The same WatsonxChatModel.chat task might output either a natural language response or a structured tool call like get_weather(latitude=40, longitude=-74) .

3. **General attributes:** Provide telemetry metadata: token usage, timing information, identifiers like unique IDs and model information.

Example: You might see that a task used 450 input tokens and 120 output tokens, took 1.2 seconds to execute and used the ibm/granite-3.1-8b-instruct model .

How to use task attributes

Together, these attributes let you fully understand what the model saw, what it decided and how it responded.

This level of detail is invaluable for debugging, optimization and validation.

Understanding task metrics

Every task includes performance and cost-related metrics that summarize how the task executed. These metrics provide quantitative data about agent performance.

Key metrics include:

- **Total execution time:** How long the task took from start to finish
- **LLM call count:** How many times the language model was invoked
- **Tool call count:** How many times external tools were called
- **Token usage:** Input tokens, output tokens and total tokens consumed
- **Cost estimates:** Approximate costs based on token usage (when pricing data is available)
- **Subtask distribution:** How work was distributed across child tasks

These metrics help you optimize performance and debug agent behavior. They can also aid in identifying slow tasks that could be parallelized or cached. This view is integral for capacity planning because it allows you to understand resource requirements for scaling and track token usage to control expenses.

For example, if you notice that a trace took 8 seconds but only 0.5 seconds were spent on LLM calls, you know that the bottleneck is elsewhere (likely in tool execution or network latency).

Understanding agent spans

While tasks show you the logical workflow of your agent, **spans** represent the underlying system-level operations that occur during execution. Clicking the **Spans** tab reveals what the platform is doing internally to process each request.

Spans provide visibility into the low-level execution steps recorded by the orchestration framework (in this case, LangGraph, an open source framework running inside the wxo-server). Each span represents a discrete operation such as:

- Routing the request to the correct agent (agent_style_router)
- Invoking the agent and initializing its context (agent.task)
- Building prompts and context from templates (ChatPromptTemplate.task)
- Calling the LLM with the assembled prompt (WatsonxChatModel.chat)
- Returning results to the user (answer.task)

How spans differ from tasks

While tasks show the *logical* steps of agent execution (what the agent is trying to accomplish), spans show the *technical steps* (how the system accomplishes it). This dual view gives you both the high-level understanding and low-level debugging capability.

Example: A single task like 0.1:agent might contain multiple spans representing database queries, cache lookups and configuration loading. These operations happen behind the scenes to support the agent's execution.

Understanding span tags

Each span includes **tags** that provide additional metadata and context. These tags are essential for filtering, debugging and analyzing agent performance.

Common span tags include:

- Agent identification: agent_id , agent_name
- Session tracking: thread_id , session_id , conversation_id
- Workflow context: step_number , workflow_path , parent_span_id
- Performance data: token_count , duration_ms , model_name
- Request details: tool_calls , input_preview , output_preview

Using spans for debugging

Spans are useful for tracking down latency, understanding failures by seeing which internal component failed, analyzing patterns by filtering spans by tag to identify trends and cross-referencing by linking spans across multiple traces by using session IDs.

For example, if your agent occasionally hangs, you can filter spans by duration to identify which internal operations are taking unexpectedly long, perhaps a database query or a network call to an external service.

Visualizing execution with the Workflows tab

The **Workflows** tab provides a hierarchical visualization called the **Runnables Tree**, which shows the complete execution structure of your agent workflow. This view is especially useful for understanding complex multi-agent systems and nested execution patterns.

What is a runnable?

In the Watsonx Orchestrate framework, a **runnable** is a unit of work or task that can be executed. Runnables can be:

- **Simple operations:** A single LLM call or tool invocation
- **Composite workflows:** Multiple runnables chained together
- **Conditional branches:** Different execution paths based on conditions
- **Parallel executions:** Multiple runnables running simultaneously

Understanding the tree structure

The Runnable Tree displays parent-child relationships, making it easy to see:

- Which tasks trigger others: Following the execution chain
- Parallel vs. sequential execution: Understanding workflow concurrency
- Branching logic: How decisions lead to different execution paths
- Workflow depth: How deeply nested your agent logic is

When workflows become critical

For simple agents like the Weather Agent, the workflow view mirrors the task view closely. However, workflows become invaluable when you're working with:

- **Multi-agent systems:** Multiple specialized agents collaborating on a task.
- **Complex orchestration:** Agents that dynamically choose between different tools or sub-agents
- **Iterative refinement:** Agents that loop through steps until a condition is met
- **Conditional routing:** Workflows that branch based on intermediate results
- **Scalable architectures:** Designing workflows that handle real-world loads efficiently

For example, imagine an agent that first checks if a query requires web search, then decides between using a calculator tool or a database query tool, and finally validates the result before responding. The Runnables Tree would show this entire branching structure clearly.

Using the workflow view

You can interact with the tree by:

- **Expanding/collapsing nodes:** Focus on specific workflow sections
- **Clicking nodes:** Jump to detailed task information
- **Following execution paths:** Trace how data flows through the window
- **Identifying bottlenecks:** Spot where workflows become inefficient

The visualization makes debugging workflows significantly easier than trying to follow text logs or trace data alone.

Advanced analytics: The Eval tab

The **Eval** (evaluation) tab provides a quality assurance and monitoring view that measures the correctness and reliability of your agent's execution. This step is where you move from observing *what*

happened to evaluating how well it happened.

The Eval tab shows evaluation results that assess quality through guardrails:

- **Task success:** Which tasks completed vs. failed
- **Output quality:** Whether outputs matched expected results or quality criteria
- **Performance scores:** Quantitative metrics indicating success levels
- **Error analysis:** Categorization and severity of failures
- **Use case validation:** Whether agent behavior matches intended use cases

Evaluations help you monitor reliability by tracking how consistently your agent produces correct results, identify when changes degrade agent performance, prioritize improvements and build confidence by validating that agents work correctly before production deployment.

You can leverage evaluation measurements to set up alerts, track improvements, identify patterns and use feedback to guide development to improve prompts or tools.

If you notice that 15% of weather queries fail evaluation, you can investigate those specific traces to understand whether the issue is bad input handling, API failures or incorrect response formatting.

Identifying problems with the Issues tab

The **Issues** tab provides a centralized view of everything that went wrong during workflow execution. This tab is your first stop when debugging agent failures or unexpected behavior.

The Issues tab lists problems such as:

- **Failed API calls:** External services returning errors
- **Tool execution failures:** Tools that crashed or timed out
- **Missing inputs:** Required data not available when needed
- **Model exceptions:** LLM errors like token limits or invalid inputs
- **Validation errors:** Data that doesn't meet expected formats
- **Timeout errors:** Operations that exceeded time limits
- **Unhandled runtime failures:** Unexpected exceptions in agent code

In the screenshot above, you can see a **Tool Error** that occurred when the weather API returned a 424 (Failed Dependency) or 404 (Not Found) error. The issues tab shows:

1. **The error type:** “Tool Error”
2. **The specific tool:** get_weather
3. **The error response:** Full API response showing the failure
4. **A direct link:** Click to jump to the exact task where it failed

This approach makes it simple to understand what went wrong without digging through logs or trace data.

The Issues tab is especially valuable because it aggregates failures instead of forcing you to hunt through individual tasks. It provides complete context by including full error details and related data, while severity levels enable quick triage so you can prioritize which issues to address first. The direct links to source tasks mean that one click takes you to the exact execution point where things went wrong.

Understanding agent behavior with the Trajectory tab

The **Trajectory** tab provides a chronological, conversation-style view of the agent interaction between the user and any tools the agent invokes. This view is invaluable for understanding the full context and flow of agent behavior.

The Trajectory view is useful because it allows you to see exactly how the agent processes requests from start to finish, giving you complete visibility into agent behavior. You can validate tool integration by ensuring tools are called with correct parameters and receive appropriate responses. When debugging unexpected responses, the trajectory helps you trace where the logic diverged from your expectations. You can also analyze how context builds over multiple conversation turns, watching the workflow evolve naturally. Beyond debugging, the trajectory serves as documentation, letting you capture examples of correct behavior that can be shared with team members or used as reference cases for future development. This view is particularly valuable for teams building with generative AI who need to validate agent adaptability across diverse scenarios.

Anatomy of a Trajectory

Let's walk through the Weather Agent trajectory shown in the screenshot:

1. The user query

User: "What's the weather like in NYC?"

Powered by  Granite

The conversation starts with a clear, specific request about weather in New York City.

2. Agent makes a tool call

The agent recognizes it needs external data and invokes the weather tool:

{ "current_weather": "true", "latitude": "40", "longitude": "-74" }

Powered by  Granite

This example shows that the agent correctly identified NYC's approximate coordinates, properly structured the request for the API, and set the appropriate flag for current weather.

IBM Telemetry displays this result both as raw JSON and a nicely parsed, expandable tree view.

3. The tool returns data

The weather API responds with structured weather data:

{ "temperature": "7.8", "temperature_unit": "celsius", "time": "2024-01-15T14:30:00", "weather_code": "partly_cloudy", "wind_speed": "15", "wind_speed_unit": "kmh" }

Powered by  Granite

This example shows that the tool successfully retrieved data and the response follows the expected schema and all required fields are present. Being able to inspect the raw tool response is crucial for debugging issues where the agent misinterprets tool outputs.

4. Agent summarizes the result

Finally, the agent processes the structured data and responds naturally:

Agent: "The weather in NYC is 7.8°C..."

Powered by  Granite

The agent correctly extracted the temperature and weather code and converted the structured data into natural language. The response is concise and answers the user's question.

Key Trajectory features

The Trajectory tab also supports filtering by role to view only user messages, agent messages or tool interactions. You can also expand and collapse parts of long conversations to focus on details that matter to you. For further analysis or debugging, you can export the data as JSON to jump to linked tasks from trajectory steps for corresponding details.

Conclusion

Congratulations! You've successfully set up IBM Telemetry with Watsonx Orchestrate and learned how to monitor and analyze AI agent behavior in depth. IBM Telemetry provides multiple layers of visibility to give you complete observability into how your AI agents think, decide and act. These capabilities you've explored are crucial for effective lifecycle management of agent operations in production or integrating with other agent frameworks in your environment.

If you encounter issues or have questions, check the documentation. Most common issues are covered in the [troubleshooting guide](#). You can also review GitHub issues to see whether others have experienced similar problems.

Agent monitoring through platforms like IBM Telemetry has created a robust ecosystem for AgentOps, becoming essential as autonomous agents take on more complex tasks that involve integrating SDKs, tools and external APIs. The visibility you've gained into agent behavior enables you to create more reliable, efficient and trustworthy AI systems.