

# What are hierarchical AI agents?



## Hierarchical AI agents, defined

Hierarchical agents are artificial intelligence (AI) agents that work together in a tiered [multi-agent system](#) to accomplish complex tasks. The higher-level agents in a hierarchical system handle broader responsibilities and delegate focused subtasks to lower-level agents. The [AI agents](#) in the system communicate vertically, and at times horizontally, to keep the entire system aligned.

The layered structure separates hierarchical agent systems from single-agent or flat multi-agent systems in which all the agents share the same tier. Hierarchical agent systems excel in tackling complex problems: higher-level agents handle strategy and [AI agent orchestration](#), mid-tier agents power tactical [AI decision-making](#) and lower-level agents complete tasks assigned from above.

Industry newsletter

## The latest AI trends, brought to you by experts

Get curated insights on the most important—and intriguing—AI news. Subscribe to our weekly Think newsletter. See the [IBM Privacy Statement](#).

We use your email to validate you are who you say you are, to create your IBMid, and to contact you for account related matters.

Business email

Your subscription will be delivered in English. You will find an unsubscribe link in every newsletter. You can manage your subscriptions or unsubscribe [here](#). Refer to our [IBM Privacy Statement](#) for more information.

Your subscription will be delivered in English. You will find an unsubscribe link in every newsletter. You can manage your subscriptions or unsubscribe [here](#). Refer to our [IBM Privacy Statement](#) for more information.

## Types of hierarchical agents

Most hierarchical structures have either two or three types of agents, including:

- High-level agents
- Mid-level agents
- Low-level agents

Some hierarchical [agent architectures](#) have high-level and low-level agents, while other AI systems also include a middle layer.

## High-level agents

High-level agents sit at the top level of the agentic system and can process complex tasks, formulate strategic plans, perform task decomposition and manage processes.

In contemporary agent architectures, high-level agents are often implemented with advanced [AI models](#) such as [large language models \(LLMs\)](#). In classical hierarchical [reinforcement learning](#), high-level agents instead rely on learned policies or symbolic planners.

High-level agents are responsible for the “big picture,” managing the overall system in alignment with broader goals. Interpreting a user request, forming a plan and dividing that plan into tasks are all high-level agent responsibilities.

## Mid-level agents

Mid-level agents receive directives from high-level agents. They refine and implement plans, further decompose tasks when needed and coordinate teams of low-level agents. In hierarchies without mid-level agents, high-level agents assume these duties in addition to larger-scale strategic planning.

Mid-level agents also serve as communication facilitators, converting broad goals to specific directives while consolidating reports from low-level agents and passing them upwards. High-level agents use this feedback to monitor progress and update plans as needed.

## Low-level agents

Low-level agents are the doers in the hierarchy: they are specialized agents that perform specific tasks as directed by the mid- and high-level agents above them. Low-level agents are also referred to as sub-agents.

Low-level autonomous agents are specialized for narrow tasks and might use rule-based logic, [APIs](#), learned policies, robotic controllers or other task-specific [algorithms](#). Some low-level agents, like [simple reflex agents](#), are stateless, in that they do not maintain task context or an internal representation of the environment. Higher-level agents such as [learning agents](#), by contrast, often must maintain state or context to coordinate multi-step plans and track dependencies.

Low-level agents receive directives from above, carry out their assigned tasks, then report back on the results. Mid- and high-level agents use the results from low-level agents to inform the next steps of the process in a series of interlinked dependencies.

## Features of a hierarchical agent system

Hierarchical systems are characterized by the following key features:

- Agent hierarchy
- Task decomposition

- Specialization
- Feedback-driven coordination

## Agent hierarchy

In an enterprise, C-suite executives focus on large-scale planning, middle managers convert leadership directives into operational tasks and employees handle the brunt of the actual workload. Hierarchical agent systems apply the same structure to AI-powered [agentic workflows](#).

The structure in a hierarchical multi-agent system is known as a layered control system. A clear hierarchical structure enforces job roles across the agentic ecosystem. Each tier in the hierarchy has a designated function, whether it is strategic planning, decision making, delegation and task routing or operational task execution.

## Task decomposition

Hierarchical agent systems thrive on efficient task decomposition: the process of breaking a task into smaller subtasks. The breakdown also includes any dependencies linking separate tasks. Task decomposition is found outside [machine learning](#) in many fields, such as project management and software development.

Task decomposition is a helpful strategy in general for anyone facing an imposing challenge. Converting the challenge into a sequential list of small manageable tasks can simplify a complex project.

## Specialization

Each tier in the hierarchy has a designated role—and often, each agent within a tier is also specialized for a specific task. Agent specialization is most often seen at the lower tiers, where simpler job-focused agents automate routine tasks as needed to fulfill larger-scale directives. The emphasis on specialization streamlines workflows and leads to greater efficiency.

## Feedback-driven coordination

The agents within a hierarchical system are in constant real-time communication. Some systems use asynchronous coordination, while others are strictly synchronous.

High-level agents send directives down the chain of command, while lower-level agents report back up with the results of their tasks. High-level agents use these results to optimize the strategy going forward, creating new tasks and altering task priority to best suit the long-term goal.

Meanwhile, agents can also communicate horizontally as needed to coordinate work and share results.

## Hierarchical AI agent use cases

Hierarchical agent systems are ideal for [business process automation](#) and other real-world use cases that involve complex operational processes, such as:

- Supply chain and logistics management
- Manufacturing and production
- Agentic RAG
- Cybersecurity
- Autonomous vehicles and delivery bots

## Supply chain and logistics management

In a supply chain and logistics implementation, high-level agents would plan the overall inventory management strategy. Mid-level agents would break the strategy into localized implementations, while low-level agents would handle warehouse operations and shipping routes.

## Manufacturing and production

For manufacturing, high-level agents would craft a strategy to help ensure that production meets demand in a timely manner. Mid-level agents would convert those directives into tasks for lower-level agents to implement at a local level, such as by controlling machines and production lines in a factory.

## Agentic RAG

Agentic retrieval-augmented generation (RAG) brings agentic planning to RAG systems. A high-level manager agent can coordinate the workflow, while specialized agents handle retrieval, reranking or scoring of documents, summarization and content generation. In many agentic RAG systems, user interaction is held through a chatbot interface, though RAG can also power non-chat applications.

## Cybersecurity

Multi-agent systems in cybersecurity can use a team of agents to monitor systems and combat threats. High-level agents would coordinate long-term strategy and policy updates, while low-level agents handle rapid anomaly detection and automated containment responses.

## Autonomous vehicles and delivery bots

Fleets of autonomous vehicles might be managed by hierarchical agent systems in a manner similar to a logistics fleet. High-level agents would optimize the delivery bot inventory, considering peak order times and other relevant factors. Mid-level agents coordinate deliveries within an assigned area, while low-level agents pilot individual delivery bots.

## Hierarchical agent benefits

The benefits of a hierarchical agent system include:

- Modularity
- Efficiency
- Scalability
- Fault tolerance

## Modularity

Within a hierarchical multi-agent system, each agent or group of agents is a discrete unit. Testing, development and implementation can be handled on a per-unit basis. The modular structure allows for development, updates and debugging without impacting the entire system.

## Efficiency

Hierarchy brings clarity. Each agent in the system has a defined role. Without enforced roles, agents might fail to coordinate effectively, leading to suboptimal outcomes such as bottlenecks, higher costs, production delays and a more chaotic workflow.

## Scalability

The modularity of a hierarchical agent system enhances its scalability. Organizations can add more agents as needed to a specific group, tier or role without having to rebuild the overall system from scratch. A high-level agent can coordinate growing numbers of mid-level and low-level agents without needing significant reprogramming.

## Fault tolerance

Because hierarchical systems are modular, organizations can add redundancy within key tiers, improving overall fault tolerance. With multiple agents operating in parallel, a multi-agent system can maintain operations if one or more agents fail. Higher-level manager agents can reassign tasks and restructure workflows as needed.

## Hierarchical agent challenges

The challenges of designing and implementing a hierarchical agent system include:

- Complexity
- Rigidity
- Communication bottlenecks

## Complexity

Hierarchical systems are complex in design and require a nuanced understanding of the process or operation the system is meant to manage. When built well, a hierarchical agent system can lead to smooth automation. But when poorly designed, the downsides can include inefficient operation, conflicting agent directives and excessive costs.

## Rigidity

Highly hierarchical systems with long decision chains can struggle in high-volatility environments where cascading delays lead to outdated decisions and actions. More decentralized multi-agent architectures might be better suited for dynamic scenarios.

Enterprises wanting to automate portions of their business must be judicious when choosing which areas are stable enough to entrust to a team of AI agents. Unpredictable environments might be better left to more adaptive, quick-thinking teams of human workers.

## Communication bottlenecks

A hierarchical multi-agent system is only as good as its communication bandwidth or latency constraints. Agents need to be able to send data up and down the chain of command for the system to function smoothly.

Communication delays can result in high-level agents receiving data that is no longer reflective of real-world circumstances. Communication bottlenecks are especially acute when high-level agents rely on computationally expensive reasoning, because delays can lead to lower-level agents operating on outdated instructions.